

Link layer, LANs: roadmap

- introduction
- error detection, correction
- multiple access protocols
- **LANs**
 - **addressing, ARP**
 - Ethernet
 - switches
 - VLANs
- link virtualization: MPLS
- data center networking



- a day in the life of a web request

Link Layer: 6-47

47

MAC addresses

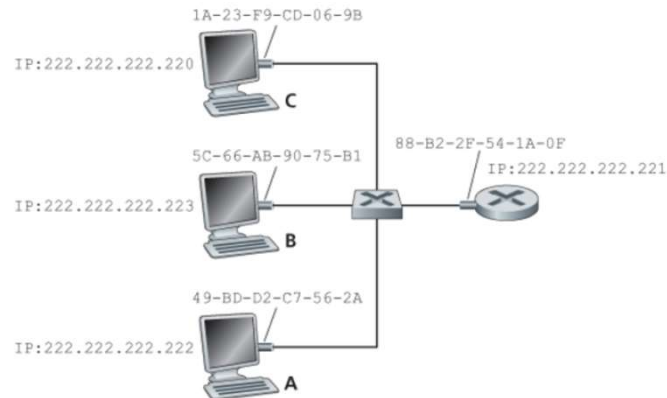
- **32-bit IP address:**
 - *network-layer* address for interface
 - used for layer 3 (network layer) forwarding
 - e.g.: **128.119.40.136**
- **MAC (or LAN or physical or Ethernet) address:**
 - function: used “locally” to get frame from one interface to another physically-connected interface (same subnet, in IP-addressing sense)
 - **48-bit** MAC address (for most LANs) burned in **NIC ROM**, also sometimes **software settable**
 - e.g.: **1A-2F-BB-76-09-AD**

*hexadecimal (base 16) notation
(each “numeral” represents 4 bits)*

Link Layer: 6-48

48

MAC addresses




Link-layer switches do not have link-layer addresses associated with their interfaces that connect to hosts and routers.

49

MAC addresses

each interface on LAN

- has unique **48-bit MAC address**
- has a locally unique 32-bit IP address (as we've seen)



Link Layer: 6-50

50

MAC addresses

- **MAC address** allocation administered by **IEEE**
- manufacturer buys **portion of MAC address space** (to assure uniqueness)
- analogy:
 - **MAC address**: like **Social Security Number**
 - **IP address**: like **postal address**
- **MAC flat address**: **portability**
 - can **move interface** from one LAN to another
 - recall **IP address *not* portable**: depends on **IP subnet** to which node is attached

IEEE : Institute of Electrical and Electronic Engineers


Link Layer: 6-51

51

ARP: address resolution protocol : RFC 826

Question: how to determine interface's MAC address, knowing its **IP address**?

ARP table: each IP node (host, router) on LAN has **table**



- IP/MAC address mappings for some LAN nodes:
< IP address; MAC address; TTL>
- **TTL (Time To Live)**: time after which address mapping will be forgotten (typically 20 min)

IP Address	MAC Address	TTL
222.222.222.221	88-B2-2F-54-1A-0F	13:45:00
222.222.222.223	5C-66-AB-90-75-B1	13:52:00

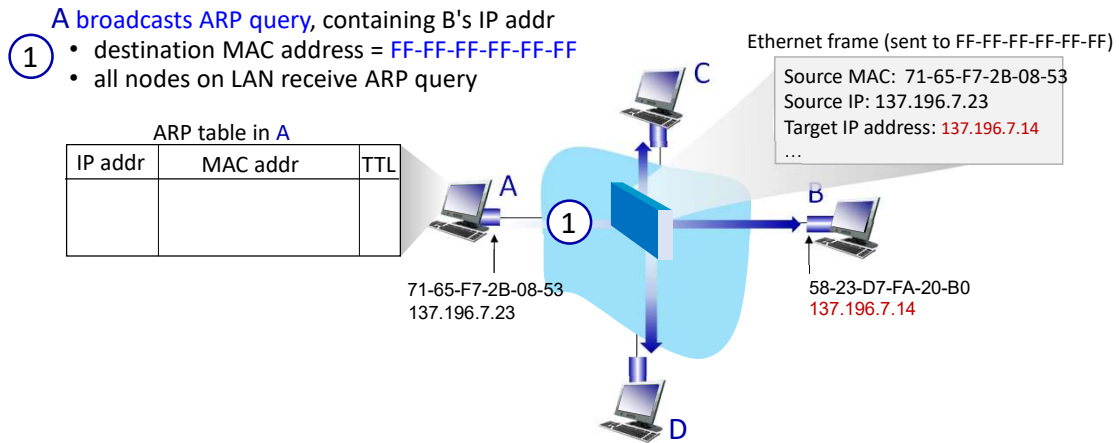
Link Layer: 6-52

52

ARP protocol in action

example: A wants to send datagram to B

- B's MAC address **not** in A's ARP table, so A uses ARP to find B's MAC address

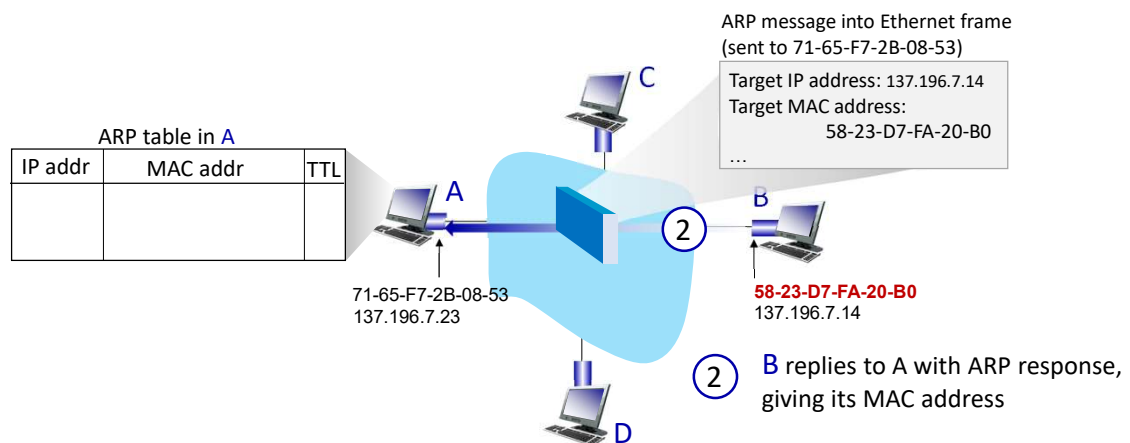


53

ARP protocol in action

example: A wants to send datagram to B

- B's MAC address not in A's ARP table, so A uses ARP to find B's MAC address

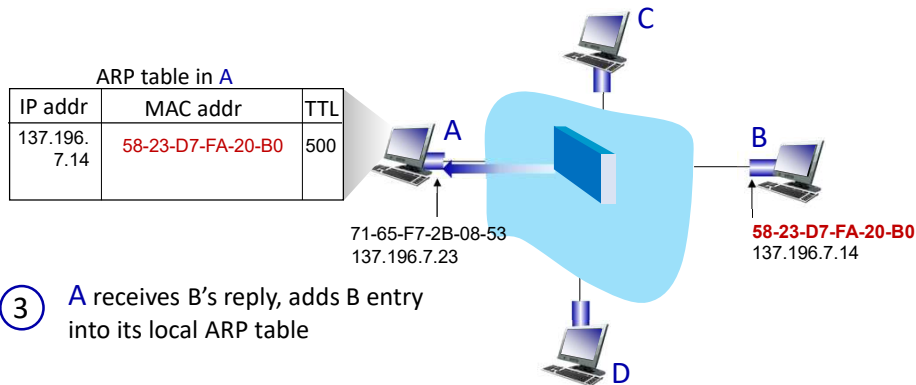


54

ARP protocol in action

example: A wants to send datagram to B

- B's MAC address **not** in A's ARP table, so A uses ARP to find B's MAC address



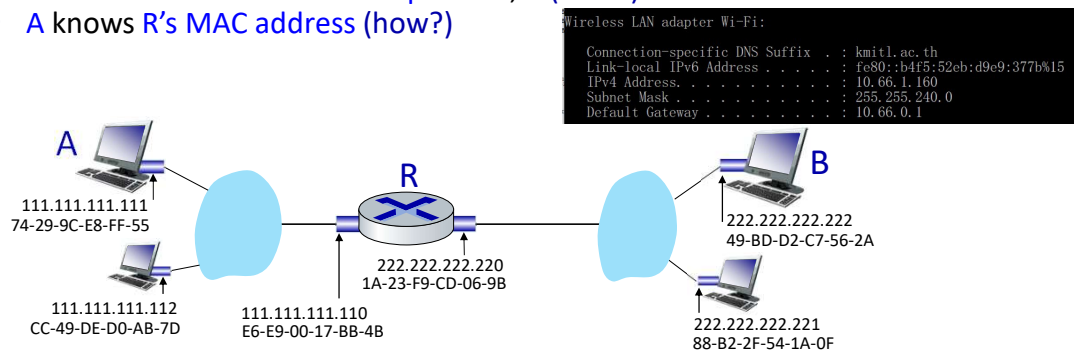
Link Layer: 6-55

55

Routing to another subnet: addressing

walkthrough: **sending a datagram from A to B via R**

- focus on addressing – at IP (datagram) and MAC layer (frame) levels
- assume that:
 - A knows B's IP address
 - A knows IP address of first hop router, R (how?)
 - A knows R's MAC address (how?)




Link Layer: 6-56

56

Routing to another subnet: addressing

- A creates IP datagram with IP source A, destination B
- A creates link-layer frame containing A-to-B IP datagram
 - R's MAC address is frame's destination




Link Layer: 6-57

57

Routing to another subnet: addressing

- frame sent from A to R
- frame received at R, datagram removed, passed up to IP




Link Layer: 6-58

58

Routing to another subnet: addressing

- R determines outgoing interface, passes datagram with IP source A, destination B to link layer
- R creates link-layer frame containing A-to-B IP datagram. Frame destination address: B's MAC address




Link Layer: 6-59

59

Routing to another subnet: addressing

- R determines outgoing interface, passes datagram with IP source A, destination B to link layer
- R creates link-layer frame containing A-to-B IP datagram. Frame destination address: B's MAC address
- transmits link-layer frame




Link Layer: 6-60

60

Routing to another subnet: addressing

- B receives frame, extracts IP datagram destination B
- B passes datagram up protocol stack to IP



Link Layer: 6-61

61

Link layer, LANs: roadmap

- introduction
- error detection, correction
- multiple access protocols
- **LANs**
 - addressing, ARP
 - **Ethernet**
 - switches
 - VLANs
- link virtualization: MPLS
- data center networking



- a day in the life of a web request


Link Layer: 6-62

62

Ethernet

“dominant” wired LAN technology:

- first widely used LAN technology
- simpler, cheap
- kept up with speed race: 10 Mbps – 400 Gbps
- single chip, multiple speeds (e.g., Broadcom BCM5761)



Metcalfe's Ethernet sketch

<https://www.uspto.gov/learning-and-resources/journeys-innovation/audio-stories/defying-doubters>


Link Layer: 6-63

63

Ethernet: physical topology

- **bus**: popular through mid 90s
 - all nodes in same collision domain (can collide with each other)
- **switched**: prevails today
 - active link-layer 2 **switch** in center
 - each “spoke” runs a (separate) Ethernet protocol (nodes do not collide with each other)

bus: coaxial cable



switched

Link Layer: 6-64

64

Ethernet frame structure

sending interface encapsulates IP datagram (or other network layer protocol packet) in **Ethernet frame**




preamble:

- used to **synchronize receiver**, sender clock rates
- **7 bytes** of **10101010** followed by one byte of **10101011**

Link Layer: 6-65

65

Ethernet frame structure (more)



- **addresses:** 6 byte source, destination MAC addresses
 - if **adapter receives frame** with **matching destination address**, or **with broadcast address** (e.g., ARP packet), it **passes data in frame** to **network layer protocol**
 - otherwise, adapter **discards frame**
- **type:** indicates higher layer protocol
 - **mostly IP** but others possible, e.g., Novell IPX, AppleTalk
 - used to demultiplex up at receiver
- **CRC:** cyclic redundancy check at receiver
 - error detected: frame is dropped

Link Layer: 6-66

66

EtherType (hexadecimal)	Protocol
0x0800	Internet Protocol version 4 (IPv4)
0x0806	Address Resolution Protocol (ARP)
0x0842	Wake-on-LAN ^[9]
0x22F0	Audio Video Transport Protocol (AVTP)
0x22F3	IETF TRILL Protocol
0x22EA	Stream Reservation Protocol
0x6002	DEC MOP RC
0x6003	DECnet Phase IV, DNA Routing
0x6004	DEC LAT
0x8035	Reverse Address Resolution Protocol (RARP)
0x809B	AppleTalk (EtherTalk)
0x80F3	AppleTalk Address Resolution Protocol (AARP)
0x8100	VLAN-tagged frame (IEEE 802.1Q) and Shortest Path Bridging IEEE 802.1aq with NNI compatibility ^[10]
0x8102	Simple Loop Prevention Protocol (SLPP)
0x8103	Virtual Link Aggregation Control Protocol (VLACP)
0x8137	IPX
0x8204	QNX Qnet
0x86DD	Internet Protocol Version 6 (IPv6)
0x8808	Ethernet flow control

Link Layer: 6-67

67

Ethernet: unreliable, connectionless


- **connectionless**: no handshaking between sending and receiving NICs
- **unreliable**: receiving NIC doesn't send ACKs or NAKs to sending NIC
 - data in dropped frames recovered only if initial sender uses higher layer rdt (e.g., TCP), otherwise dropped data lost
- Ethernet's MAC protocol: unslotted **CSMA/CD with binary backoff**

Link Layer: 6-68

68

802.3 Ethernet standards: link & physical layers

- *many* different Ethernet standards
 - common MAC protocol and frame format
 - different speeds: 10 Mbps, 100 Mbps, 1Gbps, 10 Gbps, 40 Gbps
 - different physical layer media: fiber, cable



Link Layer: 6-69

69

Link layer, LANs: roadmap

- introduction
- error detection, correction
- multiple access protocols
- **LANs**
 - addressing, ARP
 - Ethernet
 - **switches**
 - VLANs
- link virtualization: MPLS
- data center networking



- a day in the life of a web request

Link Layer: 6-71

71

Ethernet switch


- Switch is a **link-layer** device: takes an *active* role
 - store, forward Ethernet frames
 - examine incoming frame's MAC address, *selectively* forward frame to one-or-more outgoing links when frame is to be forwarded on segment, uses CSMA/CD to access segment
- transparent: hosts *unaware* of presence of switches
- plug-and-play, self-learning
 - switches do **not need** to be configured

Link Layer: 6-72

72

Switch: multiple simultaneous transmissions

- hosts have dedicated, direct connection to switch
- switches buffer packets
- Ethernet protocol used on *each* incoming link, so:
 - no collisions; full duplex
 - each link is its own collision domain
- **switching**: A-to-A' and B-to-B' can transmit simultaneously, without collisions




Link Layer: 6-73

73

Switch: multiple simultaneous transmissions

- hosts have dedicated, direct connection to switch
- switches buffer packets
- Ethernet protocol used on *each* incoming link, so:
 - no collisions; full duplex
 - each link is its own collision domain
- **switching:** A-to-A' and B-to-B' can transmit simultaneously, without collisions
 - but A-to-A' and C to A' can *not* happen simultaneously



Link Layer: 6-74

74

Switch forwarding table


Q: how does switch know A' reachable via interface 4, B' reachable via interface 5?

A: each switch has a **switch table**, each entry:

- (MAC address of host, interface to reach host, time stamp)
- looks like a routing table!

Q: how are entries created, maintained in switch table?

- something like a routing protocol?

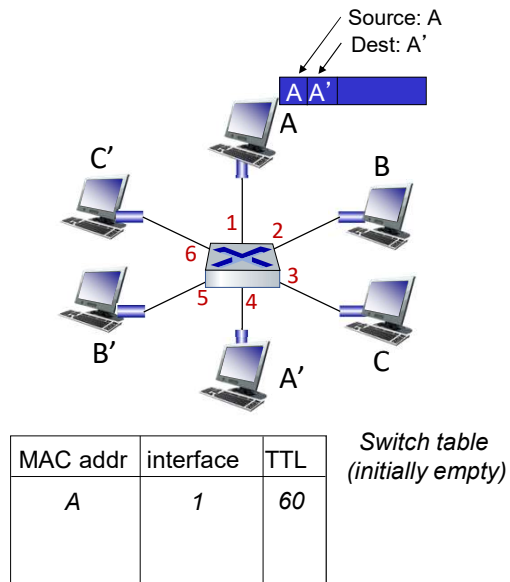


Link Layer: 6-75

75

Switch: self-learning

- switch *learns* which *hosts* can be *reached* through which *interfaces*
- when *frame received*, switch “*learns*” location of sender: *incoming* LAN segment
- records sender/*location* pair in switch table



Link Layer: 6-76

76

Switch: frame filtering/forwarding

when *frame received* at switch:


- record incoming link, MAC address of sending host
- index switch table using MAC destination address
- if entry found for destination
 - then {
 - if destination on segment from which frame arrived
 - then drop frame
 - else forward frame on interface indicated by entry
- else flood /* forward on all interfaces except arriving interface */

Link Layer: 6-77

77

Self-learning, forwarding: example

- frame destination, A', location unknown: **flood**
- destination A location known: **selectively send on just one link**



MAC addr	interface	TTL
A	1	60
A'	4	60


*switch table
(initially empty)*

Link Layer: 6-78

78

Interconnecting switches

self-learning switches can be connected together:



Q: sending from A to G - how does S₁ know to forward frame destined to G via S₄ and S₃?


- **A:** self learning! (works exactly the same as in single-switch case!)

Link Layer: 6-79

79

Self-learning multi-switch example

Suppose **C** sends **frame** to **I**, **I** responds to **C**




Q: show **switch tables** and **packet forwarding** in **S₁**, **S₂**, **S₃**, **S₄**

Link Layer: 6-80

80

Small institutional network



Link Layer: 6-81

81

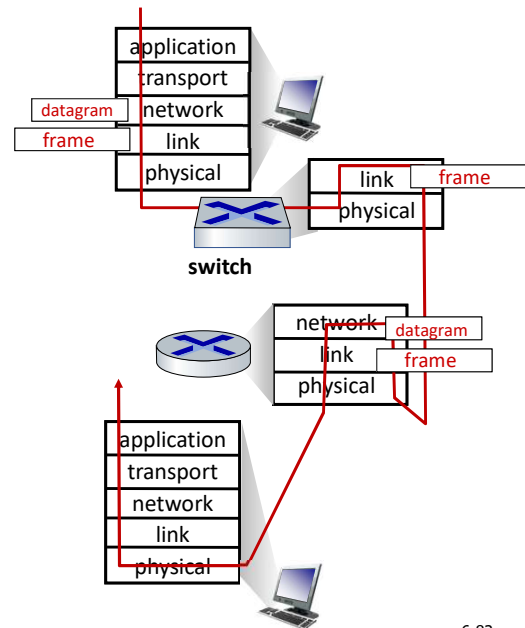
Switches vs. routers

both are store-and-forward:

- **routers:** network-layer devices (examine network-layer headers)
- **switches:** link-layer devices (examine link-layer headers)

both have forwarding tables:

- **routers:** compute tables using routing algorithms, IP addresses
- **switches:** learn forwarding table using flooding, learning, MAC addresses



Link Layer: 6-82

82

Link layer, LANs: roadmap

- introduction
- error detection, correction
- multiple access protocols
- **LANs**
 - addressing, ARP
 - Ethernet
 - switches
 - **VLANs**
- link virtualization: MPLS
- data center networking




- a day in the life of a web request

Link Layer: 6-83

83

Virtual LANs (VLANs): motivation

Q: what happens as LAN sizes scale, users change point of attachment?



single broadcast domain:


- **scaling:** all layer-2 broadcast traffic (ARP, DHCP, unknown MAC) must cross entire LAN
- efficiency, security, privacy issues

Link Layer: 6-84

84

Virtual LANs (VLANs): motivation

Q: what happens as LAN sizes scale, users change point of attachment?



single broadcast domain:

- **scaling:** all layer-2 broadcast traffic (ARP, DHCP, unknown MAC) must cross entire LAN
- efficiency, security, privacy, efficiency issues

administrative issues:

- CS user moves office to EE - **physically** attached to EE switch, but wants to remain **logically** attached to CS switch

Link Layer: 6-85


85

Port-based VLANs


Virtual Local Area Network (VLAN)

switch(es) supporting VLAN capabilities can be configured to define multiple *virtual* LANS over single physical LAN infrastructure.

port-based VLAN: switch ports grouped (by switch management software) so that *single* physical switch



... operates as **multiple** virtual switches




Link Layer: 6-86

86

Port-based VLANs


- **traffic isolation:** frames to/from ports 1-8 can *only* reach ports 1-8
 - can also define VLAN based on MAC addresses of endpoints, rather than switch port
- **dynamic membership:** ports can be dynamically assigned among VLANs
- **forwarding between VLANs:** done via routing (just as with separate switches)
 - in practice vendors sell combined switches plus routers



Link Layer: 6-87

87

VLANs spanning multiple switches



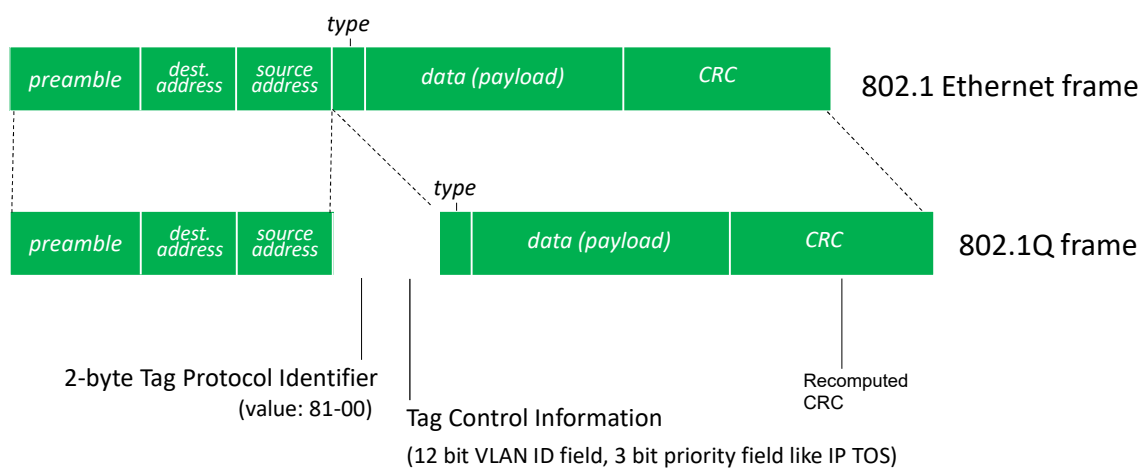
trunk port: carries frames between VLANs defined over multiple physical switches

- frames forwarded within VLAN between switches can't be vanilla 802.1 frames (must carry VLAN ID info)
- 802.1q protocol adds/removed additional header fields for frames forwarded between trunk ports

Link Layer: 6-88

88

802.1Q VLAN frame format



Link Layer: 6-89

89